

Meeting the Challenges of Solver Independence

Steven Dirkse

GAMS

Lou Hafer Simon

Fraser University

Kipp Martin

University of Chicago

Ted Ralphs


Lehigh University

November 6, 2007



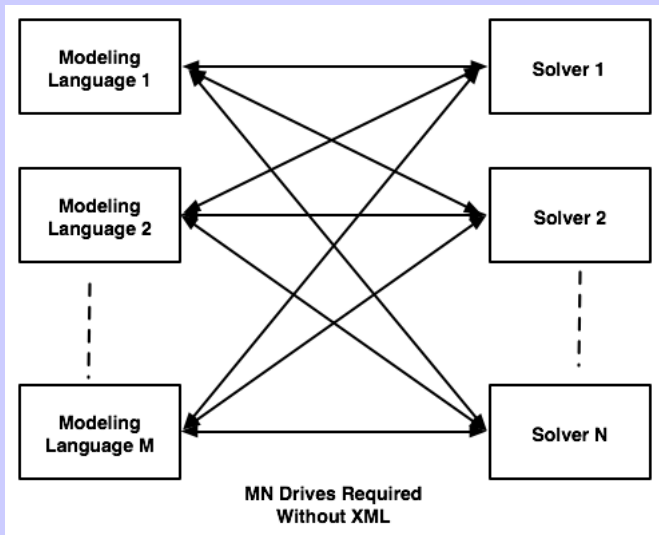
Motivation

The context: a loosely coupled modeling language and a solver.



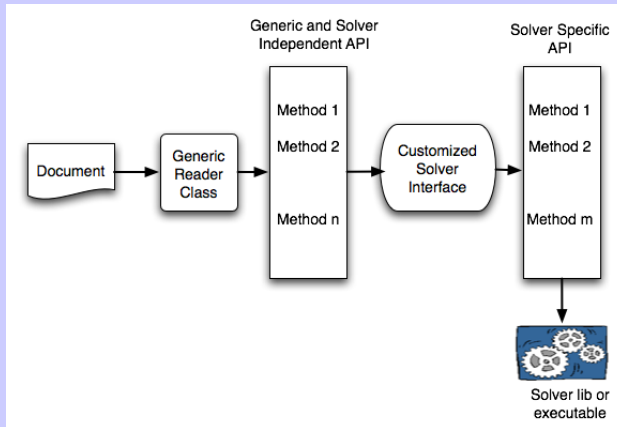
Motivation

The problem:



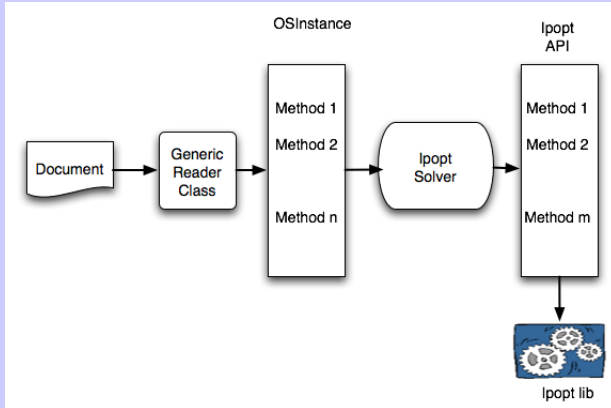
Motivation

Here is another view of solver side



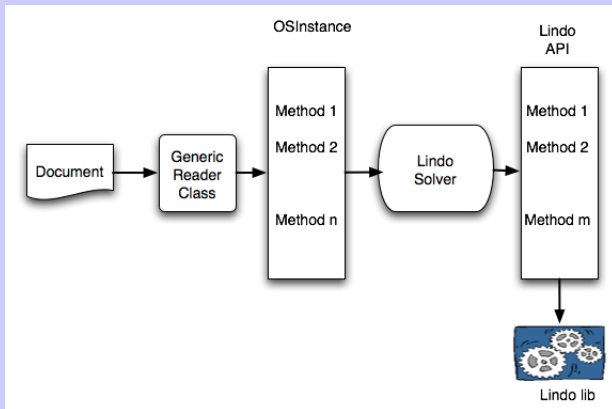
Motivation

Here is another view of solver side – with Ipopt



Motivation

Here is another view of solver side – with Lindo



Motivation

An ideal world – no solver interface!

If the instance interface and option interface are robust enough then we have:

```
load(problem_instance);  
solve(options_list);
```

It is up to the solver to implement the `problem_instance` and interpret what is in the `options_list`



COIN-OR Open Solver Interface

Problem 1: Only works for mixed-integer linear.



Extensions

- ▶ General nonlinear
- ▶ Cone
- ▶ Stochastic
- ▶ Constraint



COIN-OR Open Solver Interface

Problem 2: Does not distinguish between instance, solver options, instance result, and instance modification.



Important Concept

It is important to distinguish between:

- ▶ The model instance
- ▶ Solver options
- ▶ Solver results
- ▶ Model modifications



Important Concept

The Optimization Services project is designed to help provide this framework.

- ▶ The model instance – **OSInstance class**
- ▶ Solver options – **OSOptions class**
- ▶ Solver results – **OSResult class**
- ▶ Model modifications – to be completed



COIN-OR Open Solver Interface

Problem modification: borrow from the Gang of Four *Design Patterns*

Use the **command class** for storing modifications.



COIN-OR Open Solver Interface

Two key classes:

- ▶ OsiXFormMgr

- ▶ OsiXForm



Has pure virtual functions:

- ▶ virtual void record() = 0;
- ▶ virtual void modify() = 0;



Example

```
class ModifyConstraintBounds : public OsiXForm
```

- ▶ virtual void record() = 0;
- ▶ virtual void modify() = 0;



Algorithmic Differentiation – some motivation

Key Idea: The API of nonlinear solvers not really setup to maximize the efficient use of AD.

A typical API will have methods such as:

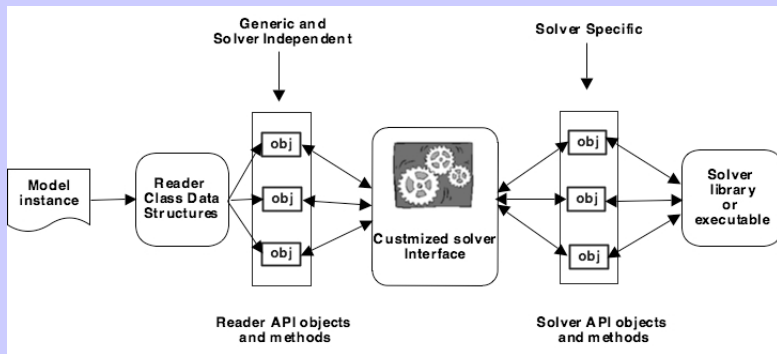
- ▶ get an objective function value
- ▶ get constraint values
- ▶ get objective gradient
- ▶ get constraint Jacobian
- ▶ get Hessian of Lagrangian

An Issue: from an AD perspective, when, for example, calculating first derivatives it would be nice to know if a second derivative is also required.



Motivation

Here is another view of solver side



Algorithmic Differentiation – some motivation

The Problem: Many nonlinear algorithms, such as interior point methods **do not** calculate all orders of derivatives for the current iterate.

For example, they may do a simple line search and not use any second derivative information.

In the API method calls for function evaluations they **do not communicate if a higher order derivative is required for the current iterate.**



Algorithmic Differentiation – gradient calculation

calculateAllConstraintFunctionGradients()

The method arguments are:

- ▶ **double* x**
- ▶ **double* objLambda**
- ▶ **double* conLambda**
- ▶ **bool new_x**
- ▶ **int highestOrder**



Algorithmic Differentiation – Hessian calculation

In our implementation, there are **exactly two** conditions that require a new function or derivative evaluation. A new evaluation is required if and only if

1. The value of `new_x` is true

–OR–

2. For the callback function the value of the input parameter `highestOrder` is strictly greater than the current value of `m_iHhighestOrderEvaluated`.

In the code we keep track of the highest order derivative calculation that has been made.

